

In []: Name: Vrushali Khade
Div: A
Roll No: 44


In [1]: `import pandas as pd`
`import numpy as np`
`import matplotlib.pyplot as plt`
`import seaborn as sns`

In [2]: `data = pd.read_csv("C:/Users/Vrushali Khade/Downloads/sales_data_sample.csv", encoding='utf-8')`
`data.head()`
While utf-8 supports all languages according to pandas' documentation, utf-8 has

Out[2]:

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	ORDERDATE
0	10107	30	95.70	2	2871.00	2010-01-01
1	10121	34	81.35	5	2765.90	2010-01-01
2	10134	41	94.74	2	3884.34	2010-01-01
3	10145	45	83.26	6	3746.70	2010-01-01
4	10159	49	100.00	14	5205.27	2010-01-01

5 rows × 7 columns



In [3]: `data.shape`

Out[3]: (2823, 7)

In [4]: *# Number of NaN values per column in the dataset*
`data.isnull().sum`

```
Out[4]: <bound method DataFrame.sum of
```

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORD
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False
...
2818	False	False	False	False
2819	False	False	False	False
2820	False	False	False	False
2821	False	False	False	False
2822	False	False	False	False

	ORDERDATE	STATUS	QTR_ID	MONTH_ID	YEAR_ID	...	ADDRESSLINE1	\
0	False	False	False	False	False	...	False	
1	False	False	False	False	False	...	False	
2	False	False	False	False	False	...	False	
3	False	False	False	False	False	...	False	
4	False	False	False	False	False	...	False	
...	
2818	False	False	False	False	False	...	False	
2819	False	False	False	False	False	...	False	
2820	False	False	False	False	False	...	False	
2821	False	False	False	False	False	...	False	
2822	False	False	False	False	False	...	False	

	ADDRESSLINE2	CITY	STATE	POSTALCODE	COUNTRY	TERRITORY	\
0	True	False	False	False	False	True	
1	True	False	True	False	False	False	
2	True	False	True	False	False	False	
3	True	False	False	False	False	True	
4	True	False	False	True	False	True	
...	
2818	True	False	True	False	False	False	
2819	True	False	True	False	False	False	
2820	True	False	True	False	False	False	
2821	True	False	True	False	False	False	
2822	True	False	False	False	False	True	

	CONTACTLASTNAME	CONTACTFIRSTNAME	DEALSIZE
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False
...
2818	False	False	False
2819	False	False	False
2820	False	False	False
2821	False	False	False
2822	False	False	False

[2823 rows x 25 columns]>

```
In [5]: data.drop(["ORDERNUMBER", "PRICEEACH", "ORDERDATE", "PHONE", "ADDRESSLINE1", "ADD
```

```
In [6]: data.head()
```

Out[6]:

	QUANTITYORDERED	ORDERLINENUMBER	SALES	STATUS	QTR_ID	MONTH_ID	YE
--	-----------------	-----------------	-------	--------	--------	----------	----

0	30	2	2871.00	Shipped	1	2	
1	34	5	2765.90	Shipped	2	5	
2	41	2	3884.34	Shipped	3	7	
3	45	6	3746.70	Shipped	3	8	
4	49	14	5205.27	Shipped	4	10	



In [7]: `data.isnull().sum()`

Out[7]:

QUANTITYORDERED	0
ORDERLINENUMBER	0
SALES	0
STATUS	0
QTR_ID	0
MONTH_ID	0
YEAR_ID	0
PRODUCTLINE	0
MSRP	0
PRODUCTCODE	0
CUSTOMERNAME	0
COUNTRY	0
DEALSIZE	0
dtype:	int64

In [8]: `data.describe()`

Out[8]:


	QUANTITYORDERED	ORDERLINENUMBER	SALES	QTR_ID	MONTH_ID
--	-----------------	-----------------	-------	--------	----------

count	2823.000000	2823.000000	2823.000000	2823.000000	2823.000000
mean	35.092809	6.466171	3553.889072	2.717676	7.09245
std	9.741443	4.225841	1841.865106	1.203878	3.65663
min	6.000000	1.000000	482.130000	1.000000	1.000000
25%	27.000000	3.000000	2203.430000	2.000000	4.000000
50%	35.000000	6.000000	3184.800000	3.000000	8.000000
75%	43.000000	9.000000	4508.000000	4.000000	11.000000
max	97.000000	18.000000	14082.800000	4.000000	12.000000



In [9]: `sns.countplot(data = data , x = 'STATUS')`


Out[9]: <Axes: xlabel='STATUS', ylabel='count'>



```
In [10]: import seaborn as sns
```

```
In [11]: sns.histplot(x = 'SALES' , hue = 'PRODUCTLINE', data = data, element="poly")
```

```
Out[11]: <Axes: xlabel='SALES', ylabel='Count'>
```



```
In [12]: data['PRODUCTLINE'].unique()
```

```
Out[12]: array(['Motorcycles', 'Classic Cars', 'Trucks and Buses', 'Vintage Cars',
               'Planes', 'Ships', 'Trains'], dtype=object)
```

```
In [13]: #checking the duplicated values
data.drop_duplicates(inplace=True)
```

```
In [14]: data.info()


<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   QUANTITYORDERED      2823 non-null   int64
 1   ORDERLINENUMBER      2823 non-null   int64
 2   SALES                 2823 non-null   float64
 3   STATUS               2823 non-null   object
 4   QTR_ID               2823 non-null   int64
 5   MONTH_ID             2823 non-null   int64
 6   YEAR_ID              2823 non-null   int64
 7   PRODUCTLINE          2823 non-null   object
 8   MSRP                 2823 non-null   int64
 9   PRODUCTCODE          2823 non-null   object
10  CUSTOMERNAME         2823 non-null   object
11  COUNTRY              2823 non-null   object
12  DEALSIZE             2823 non-null   object
dtypes: float64(1), int64(6), object(6)
memory usage: 286.8+ KB
```

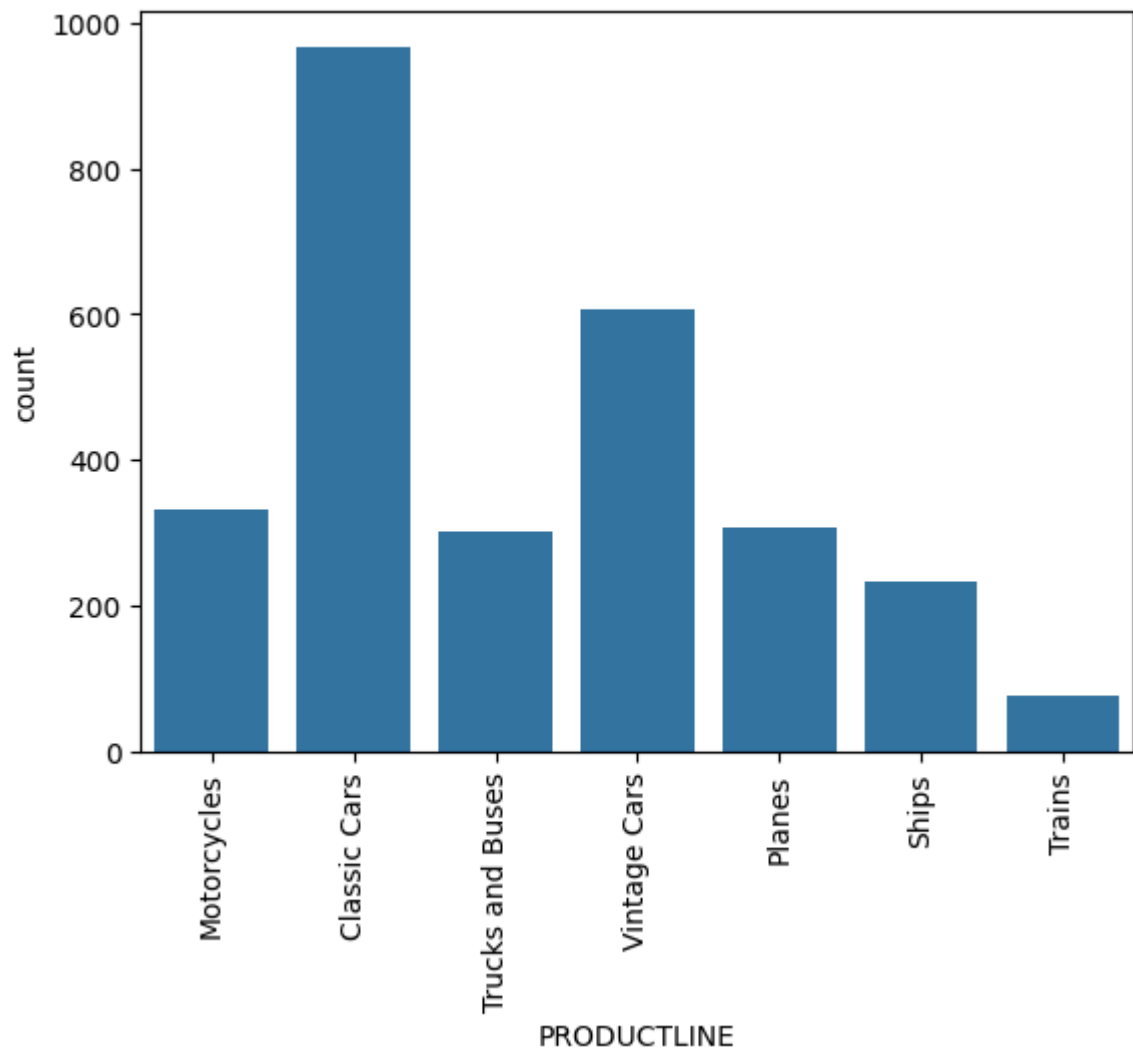
```
In [15]: list_cat = data.select_dtypes(include=['object']).columns.tolist()
```


```
In [16]: list_cat
```


```
Out[16]: ['STATUS', 'PRODUCTLINE', 'PRODUCTCODE', 'CUSTOMERNAME', 'COUNTRY', 'DEALSIZE']
```



```
In [17]: for i in list_cat:
sns.countplot(data = data ,x = i)
plt.xticks(rotation = 90)
plt.show()
```











```
In [18]: #dealing with the catagorical features
from sklearn import preprocessing
le = preprocessing.LabelEncoder()

# Encode labels in column 'species'.
for i in list_cat:
    data[i]= le.fit_transform(data[i])
```

```
In [19]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   QUANTITYORDERED        2823 non-null   int64
1   ORDERLINENUMBER        2823 non-null   int64
2   SALES                   2823 non-null   float64
3   STATUS                  2823 non-null   int32
4   QTR_ID                  2823 non-null   int64
5   MONTH_ID               2823 non-null   int64
6   YEAR_ID                 2823 non-null   int64
7   PRODUCTLINE            2823 non-null   int32
8   MSRP                    2823 non-null   int64
9   PRODUCTCODE            2823 non-null   int32
10  CUSTOMERNAME           2823 non-null   int32
11  COUNTRY                 2823 non-null   int32
12  DEALSIZE                2823 non-null   int32
dtypes: float64(1), int32(6), int64(6)
memory usage: 220.7 KB
```

```
In [20]: data['SALES'] = data['SALES'].astype(int)
```

```
In [21]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   QUANTITYORDERED        2823 non-null   int64
1   ORDERLINENUMBER        2823 non-null   int64
2   SALES                   2823 non-null   int32
3   STATUS                  2823 non-null   int32
4   QTR_ID                  2823 non-null   int64
5   MONTH_ID               2823 non-null   int64
6   YEAR_ID                 2823 non-null   int64
7   PRODUCTLINE            2823 non-null   int32
8   MSRP                    2823 non-null   int64
9   PRODUCTCODE            2823 non-null   int32
10  CUSTOMERNAME           2823 non-null   int32
11  COUNTRY                 2823 non-null   int32
12  DEALSIZE                2823 non-null   int32
dtypes: int32(7), int64(6)
memory usage: 209.6 KB
```

```
In [22]: data.describe()
```

Out[22]:

	QUANTITYORDERED	ORDERLINENUMBER	SALES	STATUS	QTR_II
count	2823.000000	2823.000000	2823.000000	2823.000000	2823.000000
mean	35.092809	6.466171	3553.421537	4.782501	2.71767
std	9.741443	4.225841	1841.865754	0.879416	1.20387
min	6.000000	1.000000	482.000000	0.000000	1.00000
25%	27.000000	3.000000	2203.000000	5.000000	2.00000
50%	35.000000	6.000000	3184.000000	5.000000	3.00000
75%	43.000000	9.000000	4508.000000	5.000000	4.00000
max	97.000000	18.000000	14082.000000	5.000000	4.00000

In [23]: *## target feature are Sales and productLine*
 X = data[['SALES', 'PRODUCTCODE']]

In [24]: data.columns

Out[24]: Index(['QUANTITYORDERED', 'ORDERLINENUMBER', 'SALES', 'STATUS', 'QTR_ID',
 'MONTH_ID', 'YEAR_ID', 'PRODUCTLINE', 'MSRP', 'PRODUCTCODE',
 'CUSTOMERNAME', 'COUNTRY', 'DEALSIZE'],
 dtype='object')

In [25]: *## K Means implementation*

In [26]: from sklearn.cluster import KMeans
 kmeans = KMeans(n_clusters=4, init='k-means++', random_state=0).fit(X)

In [27]: kmeans.labels_

Out[27]: array([2, 2, 2, ..., 3, 0, 2])

In [28]: kmeans.inertia_

Out[28]: 1043164092.8545704

In [29]: kmeans.n_iter_

Out[29]: 15

In [30]: kmeans.cluster_centers_

Out[30]: array([[1913.93425926, 63.19907407],
 [8023.78238342, 28.35751295],
 [3489.45517241, 55.50640394],
 [5371.72523364, 40.62616822]])

In [31]: *#getting the size of the clusters*
 from collections import Counter
 Counter(kmeans.labels_)

Out[31]: Counter({0: 1078, 2: 1015, 3: 537, 1: 193})

```
In [32]: sns.scatterplot(data=X, x="SALES", y="PRODUCTCODE", hue=kmeans.labels_)
plt.scatter(kmeans.cluster_centers_[0,0], kmeans.cluster_centers_[0,1],
            marker="x", c="r", s=80, label="centroids")
plt.legend()
plt.show()
```

